

# CS 636 FINAL PROJECT

## PLASTICC KAGGLE COMPETITION

### CODE

```
In [ ]: #Load libraries

import sys, os
import argparse
import time
from datetime import datetime as dt
import gc; gc.enable()
from functools import partial, wraps

import pandas as pd
import numpy as np
np.warnings.filterwarnings('ignore')

from sklearn.model_selection import StratifiedKFold
from tsfresh.feature_extraction import extract_features
from xgboost import XGBClassifier
from lightgbm import LGBMClassifier

from numba import jit
```

### Processing the metadata file and summarizing

```
In [ ]: def process_meta(filename):
    meta_df = pd.read_csv(filename)

    meta_dict = dict()
    # distance
    meta_dict.update(haversine_plus(meta_df['ra'].values, meta_df['decl'].values,
                                   meta_df['gal_l'].values, meta_df['gal_b'].values))
    #
    meta_dict['hostgal_photoz_certain'] = np.multiply(
        meta_df['hostgal_photoz'].values,
        np.exp(meta_df['hostgal_photoz_err'].values))

    meta_df = pd.concat([meta_df, pd.DataFrame(meta_dict, index=meta_df.index)], axis=1)
    return meta_df
```

## Preprocessing train\_set and test\_set

```
In [ ]: def featurize(df, df_meta, aggs, fcp, n_jobs=4):
    df = process_flux(df)

    agg_df = df.groupby('object_id').agg(aggs)
    agg_df.columns = [ '{}_{}'.format(k, agg) for k in aggs.keys() for agg in
aggs[k]]
    agg_df = process_flux_agg(agg_df)

    agg_df_ts_flux_passband = extract_features(df,
                                                column_id='object_id',
                                                column_sort='mjd',
                                                column_kind='passband',
                                                column_value='flux',
                                                default_fc_parameters=fcp['flux
_passband'], n_jobs=n_jobs)

    agg_df_ts_flux = extract_features(df,
                                      column_id='object_id',
                                      column_value='flux',
                                      default_fc_parameters=fcp['flux'], n_job
s=n_jobs)

    agg_df_ts_flux_by_flux_ratio_sq = extract_features(df,
                                                        column_id='object_id',
                                                        column_value='flux_by_flux_ratio_sq',
                                                        default_fc_parameters=fcp['flux_by_flux_
ratio_sq'], n_jobs=n_jobs)

    df_det = df[df['detected']==1].copy()
    agg_df_mjd = extract_features(df_det,
                                  column_id='object_id',
                                  column_value='mjd',
                                  default_fc_parameters=fcp['mjd'], n_jobs=n_j
obs)
    agg_df_mjd['mjd_diff_det'] = agg_df_mjd['mjd__maximum'].values - agg_df_mj
d['mjd__minimum'].values
    del agg_df_mjd['mjd__maximum'], agg_df_mjd['mjd__minimum']

    agg_df_ts_flux_passband.index.rename('object_id', inplace=True)
    agg_df_ts_flux.index.rename('object_id', inplace=True)
    agg_df_ts_flux_by_flux_ratio_sq.index.rename('object_id', inplace=True)
    agg_df_mjd.index.rename('object_id', inplace=True)
    agg_df_ts = pd.concat([agg_df,
                           agg_df_ts_flux_passband,
                           agg_df_ts_flux,
                           agg_df_ts_flux_by_flux_ratio_sq,
                           agg_df_mjd], axis=1).reset_index()

    result = agg_df_ts.merge(right=df_meta, how='left', on='object_id')
    return result
```

## Creating new fields

```
In [ ]: @jit
def haversine_plus(lon1, lat1, lon2, lat2):
    #Convert decimal degrees to Radians:
    lon1 = np.radians(lon1)
    lat1 = np.radians(lat1)
    lon2 = np.radians(lon2)
    lat2 = np.radians(lat2)

    #Implementing Haversine Formula:
    dlon = np.subtract(lon2, lon1)
    dlat = np.subtract(lat2, lat1)

    a = np.add(np.power(np.sin(np.divide(dlat, 2)), 2),
                np.multiply(np.cos(lat1),
                            np.multiply(np.cos(lat2),
                                        np.power(np.sin(np.divide(dlon, 2)), 2))))

    haversine = np.multiply(2, np.arcsin(np.sqrt(a)))
    return {
        'haversine': haversine,
        'latlon1': np.subtract(np.multiply(lon1, lat1), np.multiply(lon2, lat2
    )),
    }

@jit
def process_flux(df):
    flux_ratio_sq = np.power(df['flux'].values / df['flux_err'].values, 2.0)

    df_flux = pd.DataFrame({
        'flux_ratio_sq': flux_ratio_sq,
        'flux_by_flux_ratio_sq': df['flux'].values * flux_ratio_sq,},
        index=df.index)

    return pd.concat([df, df_flux], axis=1)

@jit
def process_flux_agg(df):
    flux_w_mean = df['flux_by_flux_ratio_sq_sum'].values / df['flux_ratio_sq_sum'].values
    flux_diff = df['flux_max'].values - df['flux_min'].values

    df_flux_agg = pd.DataFrame({
        'flux_w_mean': flux_w_mean,
        'flux_diff1': flux_diff,
        'flux_diff2': flux_diff / df['flux_mean'].values,
        'flux_diff3': flux_diff / flux_w_mean,
    }, index=df.index)

    return pd.concat([df, df_flux_agg], axis=1)
```

## Prediction functions

```

In [ ]: def predict_chunk(df_, clfs_, meta_, features, featurize_configs, train_mean):

    # process all features
    full_test = featurize(df_, meta_,
                          featurize_configs['aggs'],
                          featurize_configs['fcp'])
    full_test.fillna(0, inplace=True)

    # Make predictions
    preds_ = None
    for clf in clfs_:
        if preds_ is None:
            preds_ = clf.predict_proba(full_test[features])
        else:
            preds_ += clf.predict_proba(full_test[features])

    preds_ = preds_ / len(clfs_)

    preds_99 = np.ones(preds_.shape[0])
    for i in range(preds_.shape[1]):
        preds_99 *= (1 - preds_[i])

    preds_df_ = pd.DataFrame(preds_,
                              columns=['class_{}'.format(s) for s in clfs_[0].classes_])
    preds_df_['object_id'] = full_test['object_id']
    preds_df_['class_99'] = 0.14 * preds_99 / np.mean(preds_99)
    return preds_df_

def process_test(clfs,
                 features,
                 featurize_configs,
                 train_mean,
                 filename='predictions.csv',
                 chunks=500000):
    start = time.time()

    meta_test = process_meta('test_set_metadata.csv')

    remain_df = None
    for i_c, df in enumerate(pd.read_csv('test_set.csv', chunksize=chunks, iterator=True)):
        unique_ids = np.unique(df['object_id'])

        new_remain_df = df.loc[df['object_id'] == unique_ids[-1]].copy()
        if remain_df is None:
            df = df.loc[df['object_id'].isin(unique_ids[:-1])]
        else:
            df = pd.concat([remain_df, df.loc[df['object_id'].isin(unique_ids[:-1])]], axis=0)
        remain_df = new_remain_df

    preds_df = predict_chunk(df_=df,
                              clfs_=clfs,
                              meta_=meta_test,

```

```
features=features,
featurize_configs=featurize_configs,
train_mean=train_mean)

if i_c == 0:
    preds_df.to_csv(filename, header=True, mode='a', index=False)
else:
    preds_df.to_csv(filename, header=False, mode='a', index=False)

del preds_df
gc.collect()
print('{:15d} done in {:5.1f} minutes' .format(
    chunks * (i_c + 1), (time.time() - start) / 60), flush=True)

preds_df = predict_chunk(df_=remain_df,
                        clfs_=clfs,
                        meta_=meta_test,
                        features=features,
                        featurize_configs=featurize_configs,
                        train_mean=train_mean)

preds_df.to_csv(filename, header=False, mode='a', index=False)
return
```

## Modelling

```

In [ ]: def multi_weighted_logloss(y_true, y_preds, classes, class_weights):
    y_p = y_preds.reshape(y_true.shape[0], len(classes), order='F')
    y_ohe = pd.get_dummies(y_true)
    y_p = np.clip(a=y_p, a_min=1e-15, a_max=1 - 1e-15)
    y_p_log = np.log(y_p)
    y_log_ones = np.sum(y_ohe.values * y_p_log, axis=0)
    nb_pos = y_ohe.sum(axis=0).values.astype(float)
    class_arr = np.array([class_weights[k] for k in sorted(class_weights.keys
    ())])
    y_w = y_log_ones * class_arr / nb_pos

    loss = - np.sum(y_w) / np.sum(class_arr)
    return loss

def lgbm_multi_weighted_logloss(y_true, y_preds):
    classes = [6, 15, 16, 42, 52, 53, 62, 64, 65, 67, 88, 90, 92, 95]
    class_weights = {6: 1, 15: 2, 16: 1, 42: 1, 52: 1, 53: 1, 62: 1, 64: 2, 65
    : 1, 67: 1, 88: 1, 90: 1, 92: 1, 95: 1}

    loss = multi_weighted_logloss(y_true, y_preds, classes, class_weights)
    return 'wloss', loss, False

def xgb_multi_weighted_logloss(y_predicted, y_true, classes, class_weights):
    loss = multi_weighted_logloss(y_true.get_label(), y_predicted,
    classes, class_weights)

    return 'wloss', loss

def save_importances(importances_):
    mean_gain = importances_[['gain', 'feature']].groupby('feature').mean()
    importances_['mean_gain'] = importances_['feature'].map(mean_gain['gain'])
    return importances_

def xgb_modeling_cross_validation(params,
    full_train,
    y,
    classes,
    class_weights,
    nr_fold=5,
    random_state=1):

    # Compute weights
    w = y.value_counts()
    weights = {i : np.sum(w) / w[i] for i in w.index}

    # Loss function
    func_loss = partial(xgb_multi_weighted_logloss,
    classes=classes,
    class_weights=class_weights)

    clfs = []
    importances = pd.DataFrame()
    folds = StratifiedKFold(n_splits=nr_fold,
    shuffle=True,

```

```

        random_state=random_state)

oof_preds = np.zeros((len(full_train), np.unique(y).shape[0]))
for fold_, (trn_, val_) in enumerate(folds.split(y, y)):
    trn_x, trn_y = full_train.iloc[trn_], y.iloc[trn_]
    val_x, val_y = full_train.iloc[val_], y.iloc[val_]

    clf = XGBClassifier(**params)
    clf.fit(
        trn_x, trn_y,
        eval_set=[(trn_x, trn_y), (val_x, val_y)],
        eval_metric=func_loss,
        verbose=100,
        early_stopping_rounds=50,
        sample_weight=trn_y.map(weights)
    )
    clfs.append(clf)

    oof_preds[val_, :] = clf.predict_proba(val_x, ntree_limit=clf.best_ntree_limit)
    print('no {}-fold loss: {}'.format(fold_ + 1,
        multi_weighted_logloss(val_y, oof_preds[val_, :],
            classes, class_weights)))

    imp_df = pd.DataFrame({
        'feature': full_train.columns,
        'gain': clf.feature_importances_,
        'fold': [fold_ + 1] * len(full_train.columns),
    })
    importances = pd.concat([importances, imp_df], axis=0)

score = multi_weighted_logloss(y_true=y, y_preds=oof_preds,
    classes=classes, class_weights=class_weights)
s)
print('MULTI WEIGHTED LOG LOSS: {:.5f}'.format(score))
df_importances = save_importances(importances_=importances)
df_importances.to_csv('xgb_importances.csv', index=False)

return clfs, score

def lgbm_modeling_cross_validation(params,
    full_train,
    y,
    classes,
    class_weights,
    nr_fold=5,
    random_state=1):

    w = y.value_counts()
    weights = {i : np.sum(w) / w[i] for i in w.index}

    clfs = []
    importances = pd.DataFrame()
    folds = StratifiedKFold(n_splits=nr_fold,
        shuffle=True,
        random_state=random_state)

```



```

oof_preds = np.zeros((len(full_train), np.unique(y).shape[0]))
for fold_, (trn_, val_) in enumerate(folds.split(y, y)):
    trn_x, trn_y = full_train.iloc[trn_], y.iloc[trn_]
    val_x, val_y = full_train.iloc[val_], y.iloc[val_]

    clf = LGBMClassifier(**params)
    clf.fit(
        trn_x, trn_y,
        eval_set=[(trn_x, trn_y), (val_x, val_y)],
        eval_metric=lgbm_multi_weighted_logloss,
        verbose=100,
        early_stopping_rounds=50,
        sample_weight=trn_y.map(weights)
    )
    clfs.append(clf)

    oof_preds[val_, :] = clf.predict_proba(val_x, num_iteration=clf.best_i
teration_)
    print('no {}-fold loss: {}'.format(fold_ + 1,
        multi_weighted_logloss(val_y, oof_preds[val_, :],
            classes, class_weights)))

    imp_df = pd.DataFrame({
        'feature': full_train.columns,
        'gain': clf.feature_importances_,
        'fold': [fold_ + 1] * len(full_train.columns),
    })
    importances = pd.concat([importances, imp_df], axis=0)

    score = multi_weighted_logloss(y_true=y, y_preds=oof_preds,
        classes=classes, class_weights=class_weight
s)
    print('MULTI WEIGHTED LOG LOSS: {:.5f}'.format(score))
    df_importances = save_importances(importances_=importances)
    df_importances.to_csv('lgbm_importances.csv', index=False)

    return clfs, score

```

## Main code

```

In [ ]: aggs = {
    'flux': ['min', 'max', 'mean', 'median', 'std', 'skew'],
    'flux_err': ['min', 'max', 'mean', 'median', 'std', 'skew'],
    'detected': ['mean'],
    'flux_ratio_sq': ['sum', 'skew'],
    'flux_by_flux_ratio_sq': ['sum', 'skew'],
}

fcf = {
    'flux': {
        'longest_strike_above_mean': None,
        'longest_strike_below_mean': None,
        'mean_change': None,
        'mean_abs_change': None,
        'length': None,
    },
    'flux_by_flux_ratio_sq': {
        'longest_strike_above_mean': None,
        'longest_strike_below_mean': None,
    },
    'flux_passband': {
        'fft_coefficient': [
            {'coeff': 0, 'attr': 'abs'},
            {'coeff': 1, 'attr': 'abs'}
        ],
        'kurtosis': None,
        'skewness': None,
    },
    'mjd': {
        'maximum': None,
        'minimum': None,
        'mean_change': None,
        'mean_abs_change': None,
    },
}

best_params = {
    'device': 'cpu',
    'objective': 'multiclass',
    'num_class': 14,
    'boosting_type': 'gbdt',
    'n_jobs': -1,
    'max_depth': 7,
    'n_estimators': 500,
    'subsample_freq': 2,
    'subsample_for_bin': 5000,
    'min_data_per_group': 100,
    'max_cat_to_onehot': 4,
    'cat_l2': 1.0,
    'cat_smooth': 59.5,
    'max_cat_threshold': 32,
    'metric_freq': 10,
}

```

```
'verbosity': -1,
'metric': 'multi_logloss',
'xgboost_dart_mode': False,
'uniform_drop': False,
'colsample_bytree': 0.5,
'drop_rate': 0.173,
'learning_rate': 0.0267,
'max_drop': 5,
'min_child_samples': 10,
'min_child_weight': 100.0,
'min_split_gain': 0.1,
'num_leaves': 7,
'reg_alpha': 0.1,
'reg_lambda': 0.00023,
'skip_drop': 0.44,
'subsample': 0.75}

meta_train = process_meta('training_set_metadata.csv')

train = pd.read_csv('training_set.csv')
full_train = featurize(train, meta_train, aggs, fcp)

if 'target' in full_train:
    y = full_train['target']
    del full_train['target']

classes = sorted(y.unique())

class_weights = {c: 1 for c in classes}
class_weights.update({c:2 for c in [64, 15]})
print('Unique classes : {}, {}'.format(len(classes), classes))
print(class_weights)

if 'object_id' in full_train:
    oof_df = full_train[['object_id']]
    del full_train['object_id']
    del full_train['hostgal_specz']
    del full_train['ra'], full_train['decl'], full_train['gal_l'], full_train[
'gal_b']
    del full_train['ddf']

train_mean = full_train.mean(axis=0)

pd.set_option('display.max_rows', 500)
print(full_train.describe().T)

full_train.fillna(0, inplace=True)

eval_func = partial(lgbm_modeling_cross_validation,
                    full_train=full_train,
                    y=y,
                    classes=classes,
                    class_weights=class_weights,
                    nr_fold=7,
                    random_state=7)

best_params.update({'n_estimators': 1100})
```

```
clfs, score = eval_func(best_params)

filename = "kaggle_submission_final.csv"
print('save to {}'.format(filename))

process_test(clfs,
             features=full_train.columns,
             featurize_configs={'aggs': aggs, 'fcp': fcp},
             train_mean=train_mean,
             filename=filename,
             chunks=5000000)

z = pd.read_csv(filename)
print("Shape BEFORE grouping: {}".format(z.shape))
z = z.groupby('object_id').mean()
print("Shape AFTER grouping: {}".format(z.shape))
z.to_csv(filename, index=True)
```